

# RTGC

---

Real-Time

Multi-Core

Continuously Concurrent

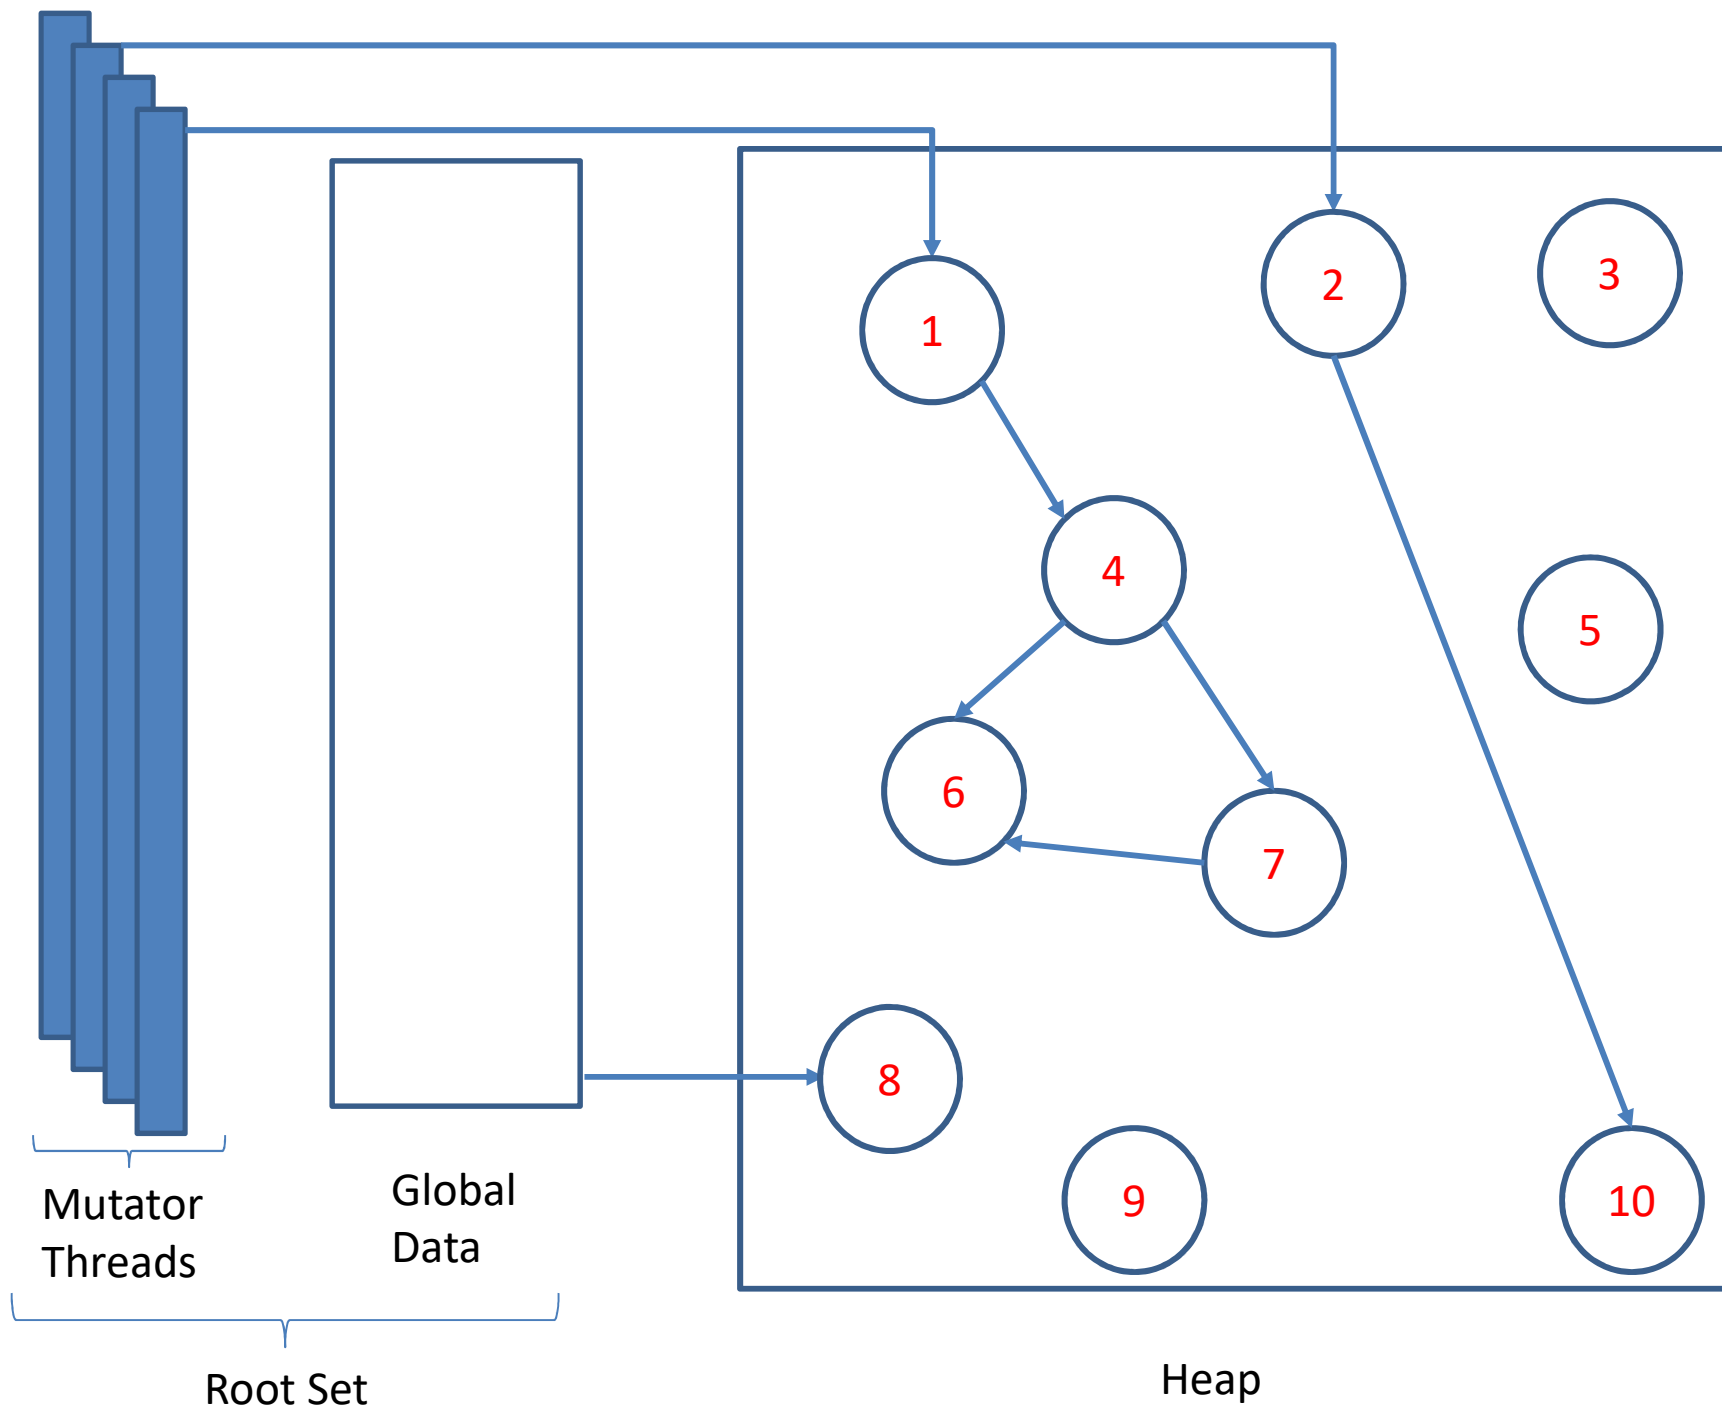
Garbage Collection

# Two Talks in One

- Part 1: Atomic GC
- Part 2: Real-time Concurrent GC

# A Visual Representation of Memory

- We need a visual representation of memory to discuss garbage collection.
- That won't fit on this slide, so it's on the next one.







# Part 1:

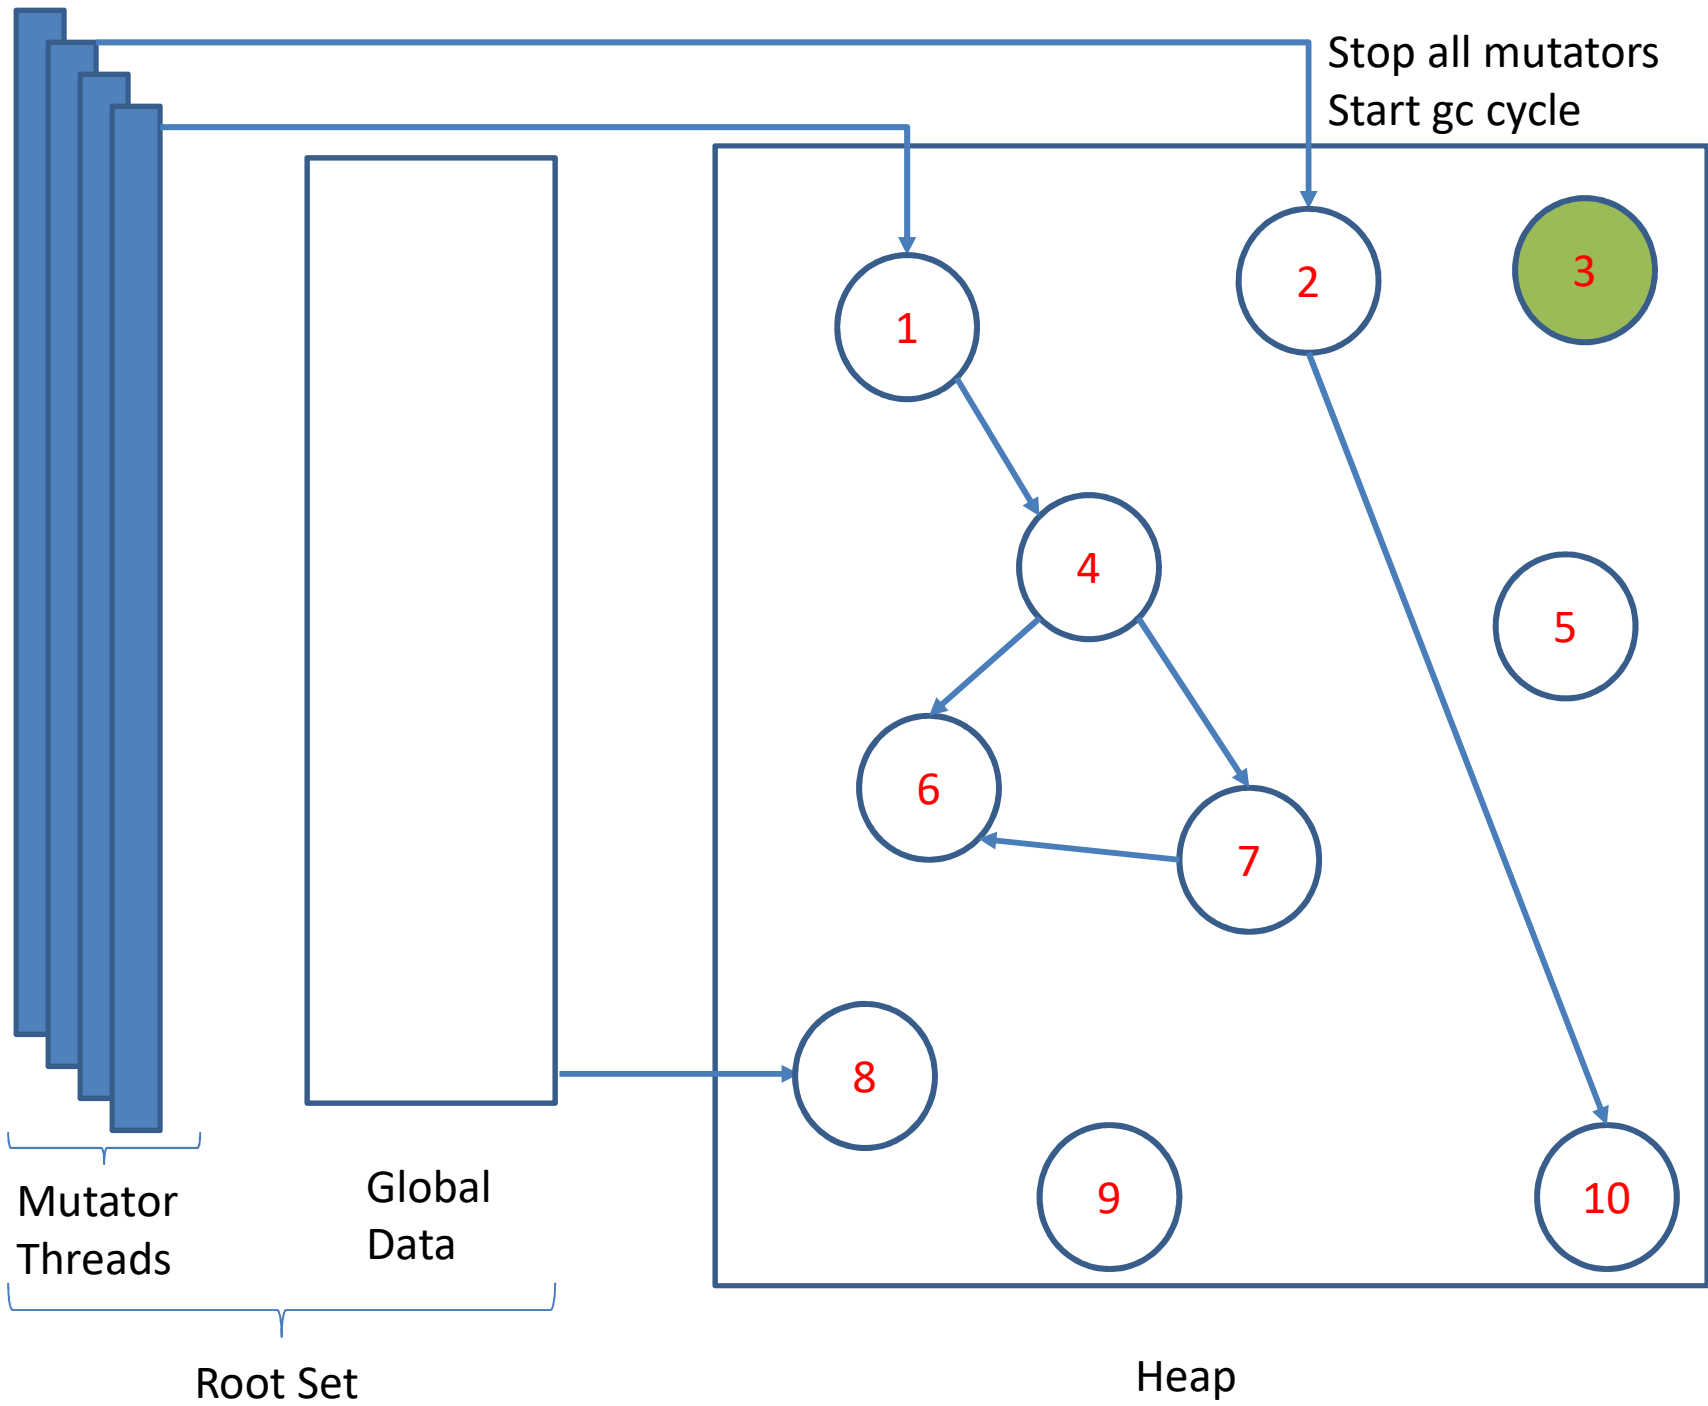
## A Simple Atomic Tracing Collector


- “Atomic” == “Stop the world” - The gc stops the execution of every mutator thread. Mutators and the gc cannot run at the same time.
- In a “gc cycle” we recursively trace the graph of all reachable objects, starting from the “root set” of pointers.
- Reclaim/collect all “garbage” objects that weren’t reachable by tracing so they may be allocated as new objects.
- No way for mutators to alter the object pointer graph while the gc works because they are all stopped.
- Resume all mutator threads after a gc cycle has completed.
- Your program is either “mutating” or “collecting”, but not both at the same time.

# Quad-Color marking


- Assign each object a color:
  -  White - All allocated objects are white at the beginning of a gc cycle. This means they may be reachable, or they may be garbage. We don't know yet.
  -  Gray - A gray object has been found to be reachable. However, we don't know anything about other objects it may point at recursively.
  -  Black - A black object is not only reachable, but we have recursively scanned it for pointers to any other objects in the graph.
  -  Green – free objects that can be allocated as new objects.
- All reachable objects change from white to gray to black in a gc cycle.
- Any objects that remain white at the end of a gc cycle are garbage and are changed to green.


Lets look at a visual example of this process.




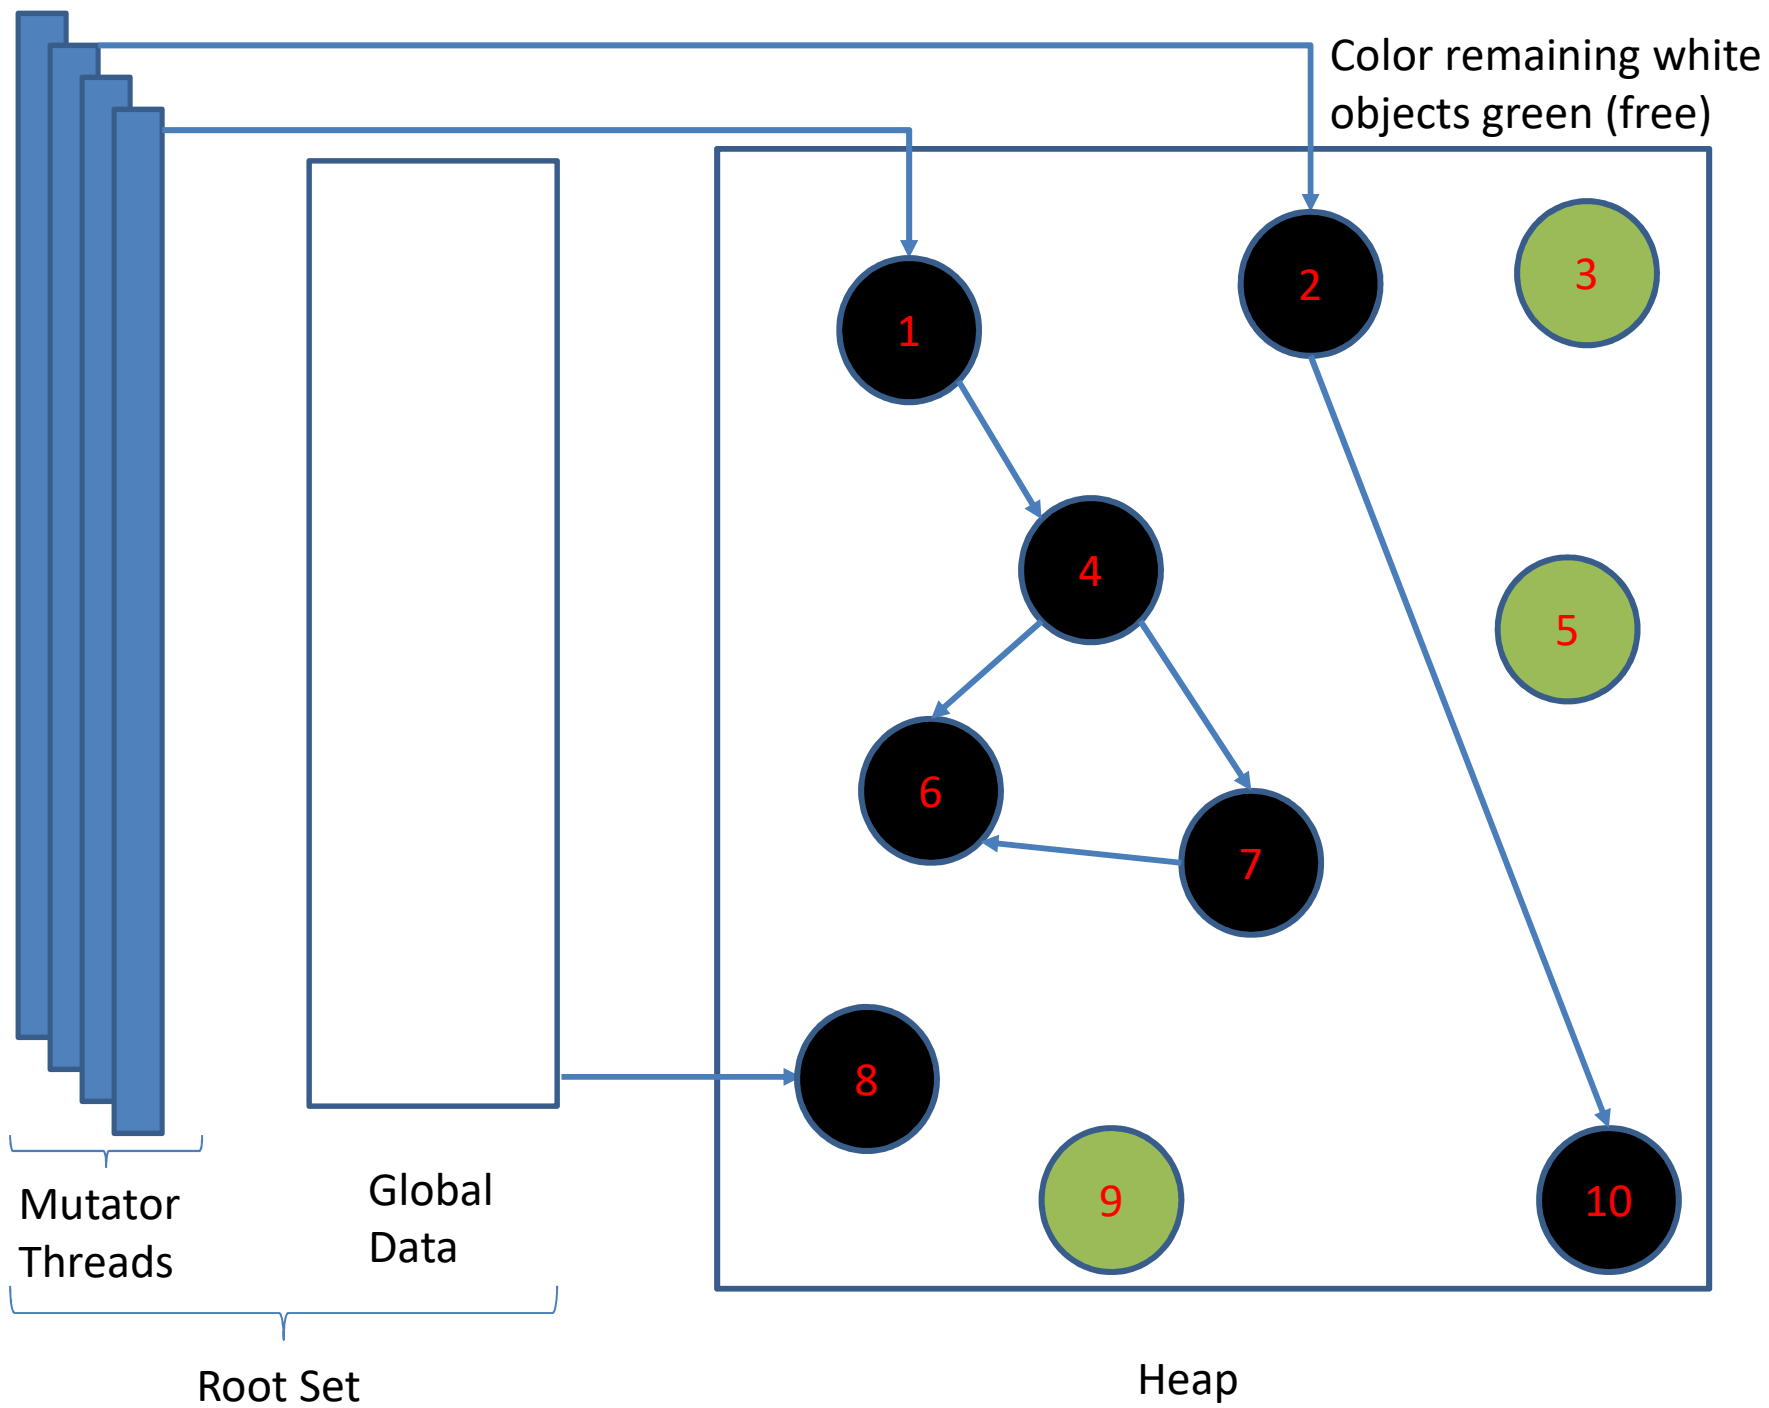









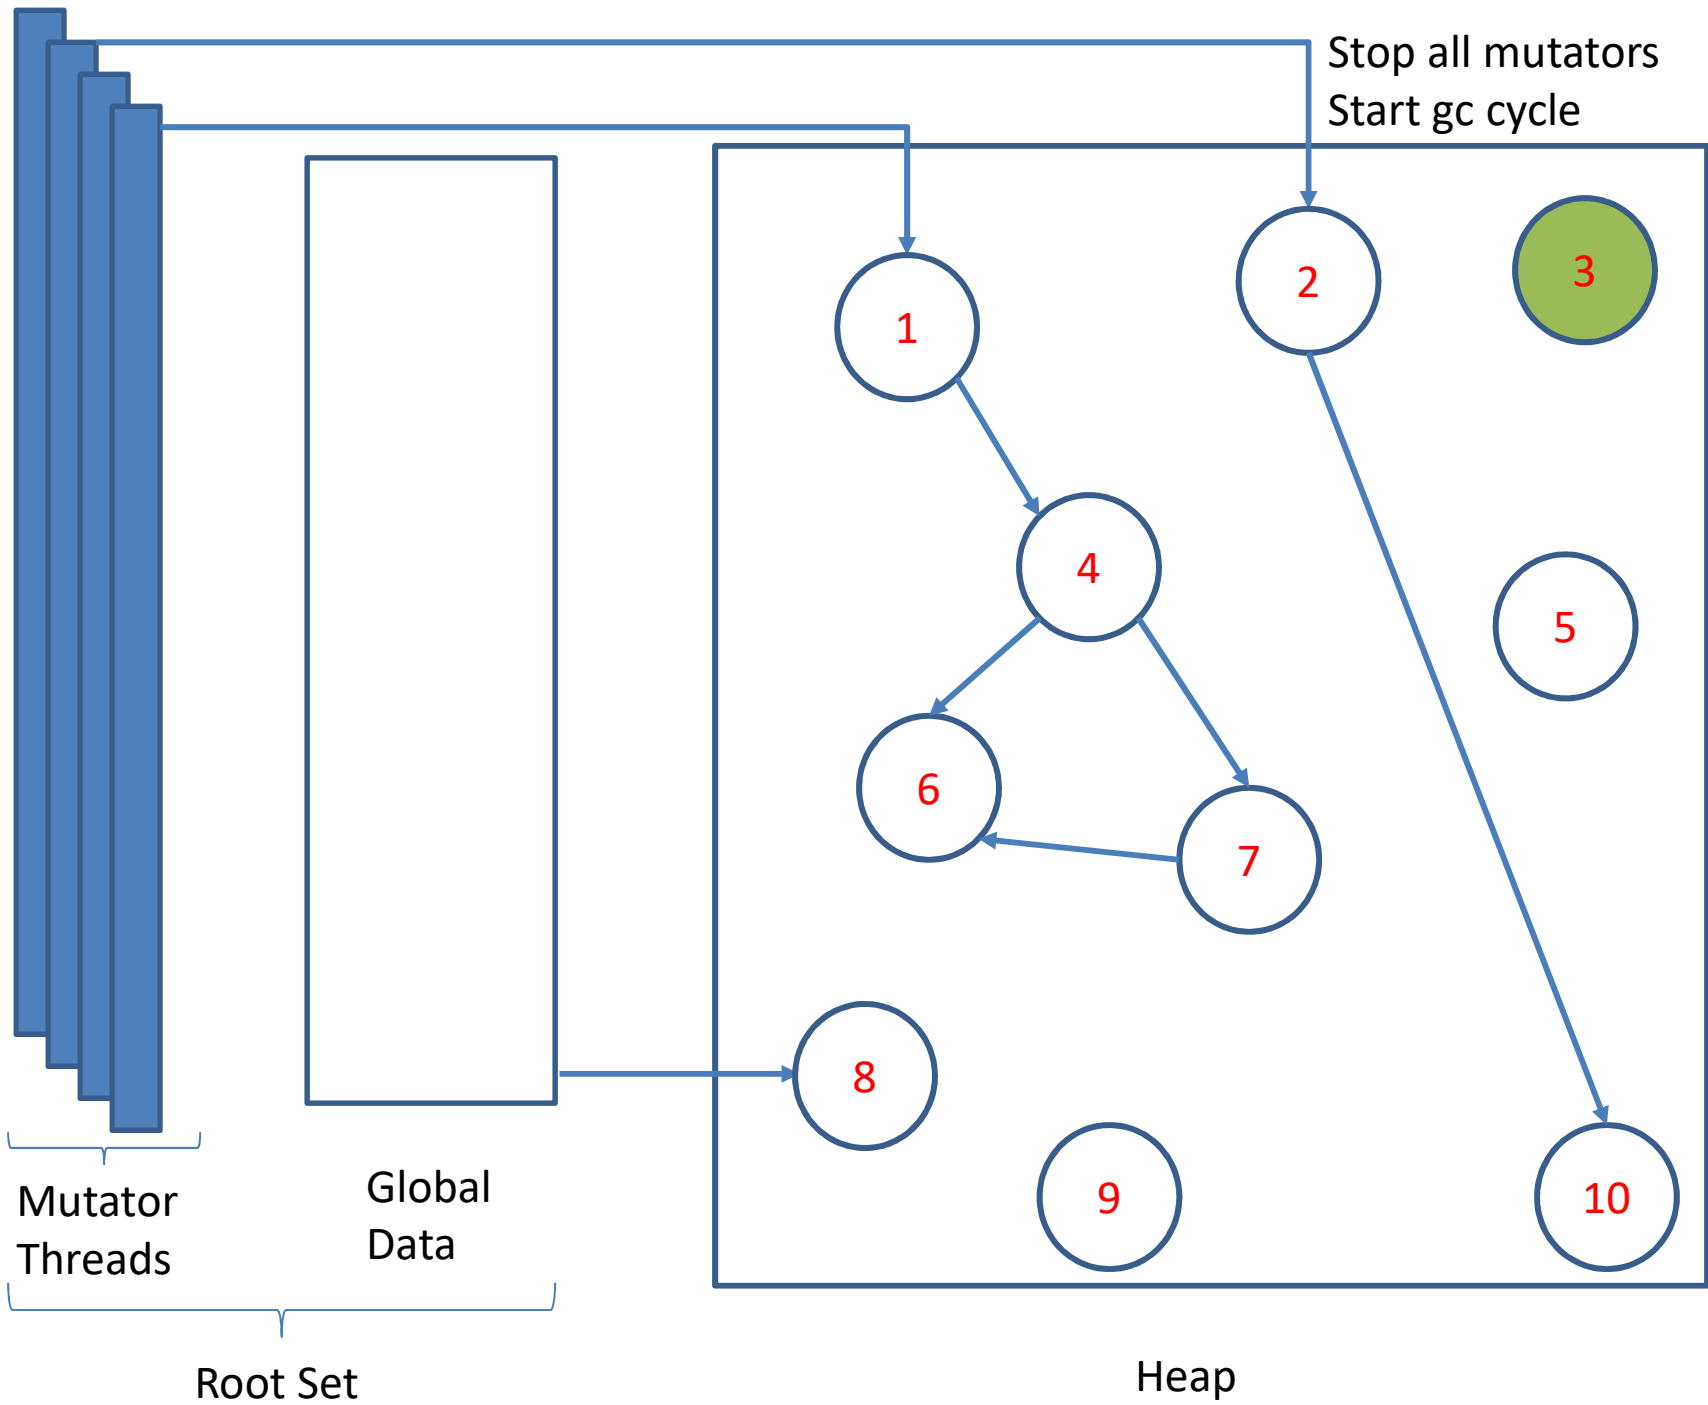







# One Problem

- We repeatedly had to look at every single object to determine which ones were gray.
- Let's fix this by treating object colors as sets that are redundantly represented by doubly-linked lists.




Black = NULL


Free Last




Gray = NULL









Black = NULL

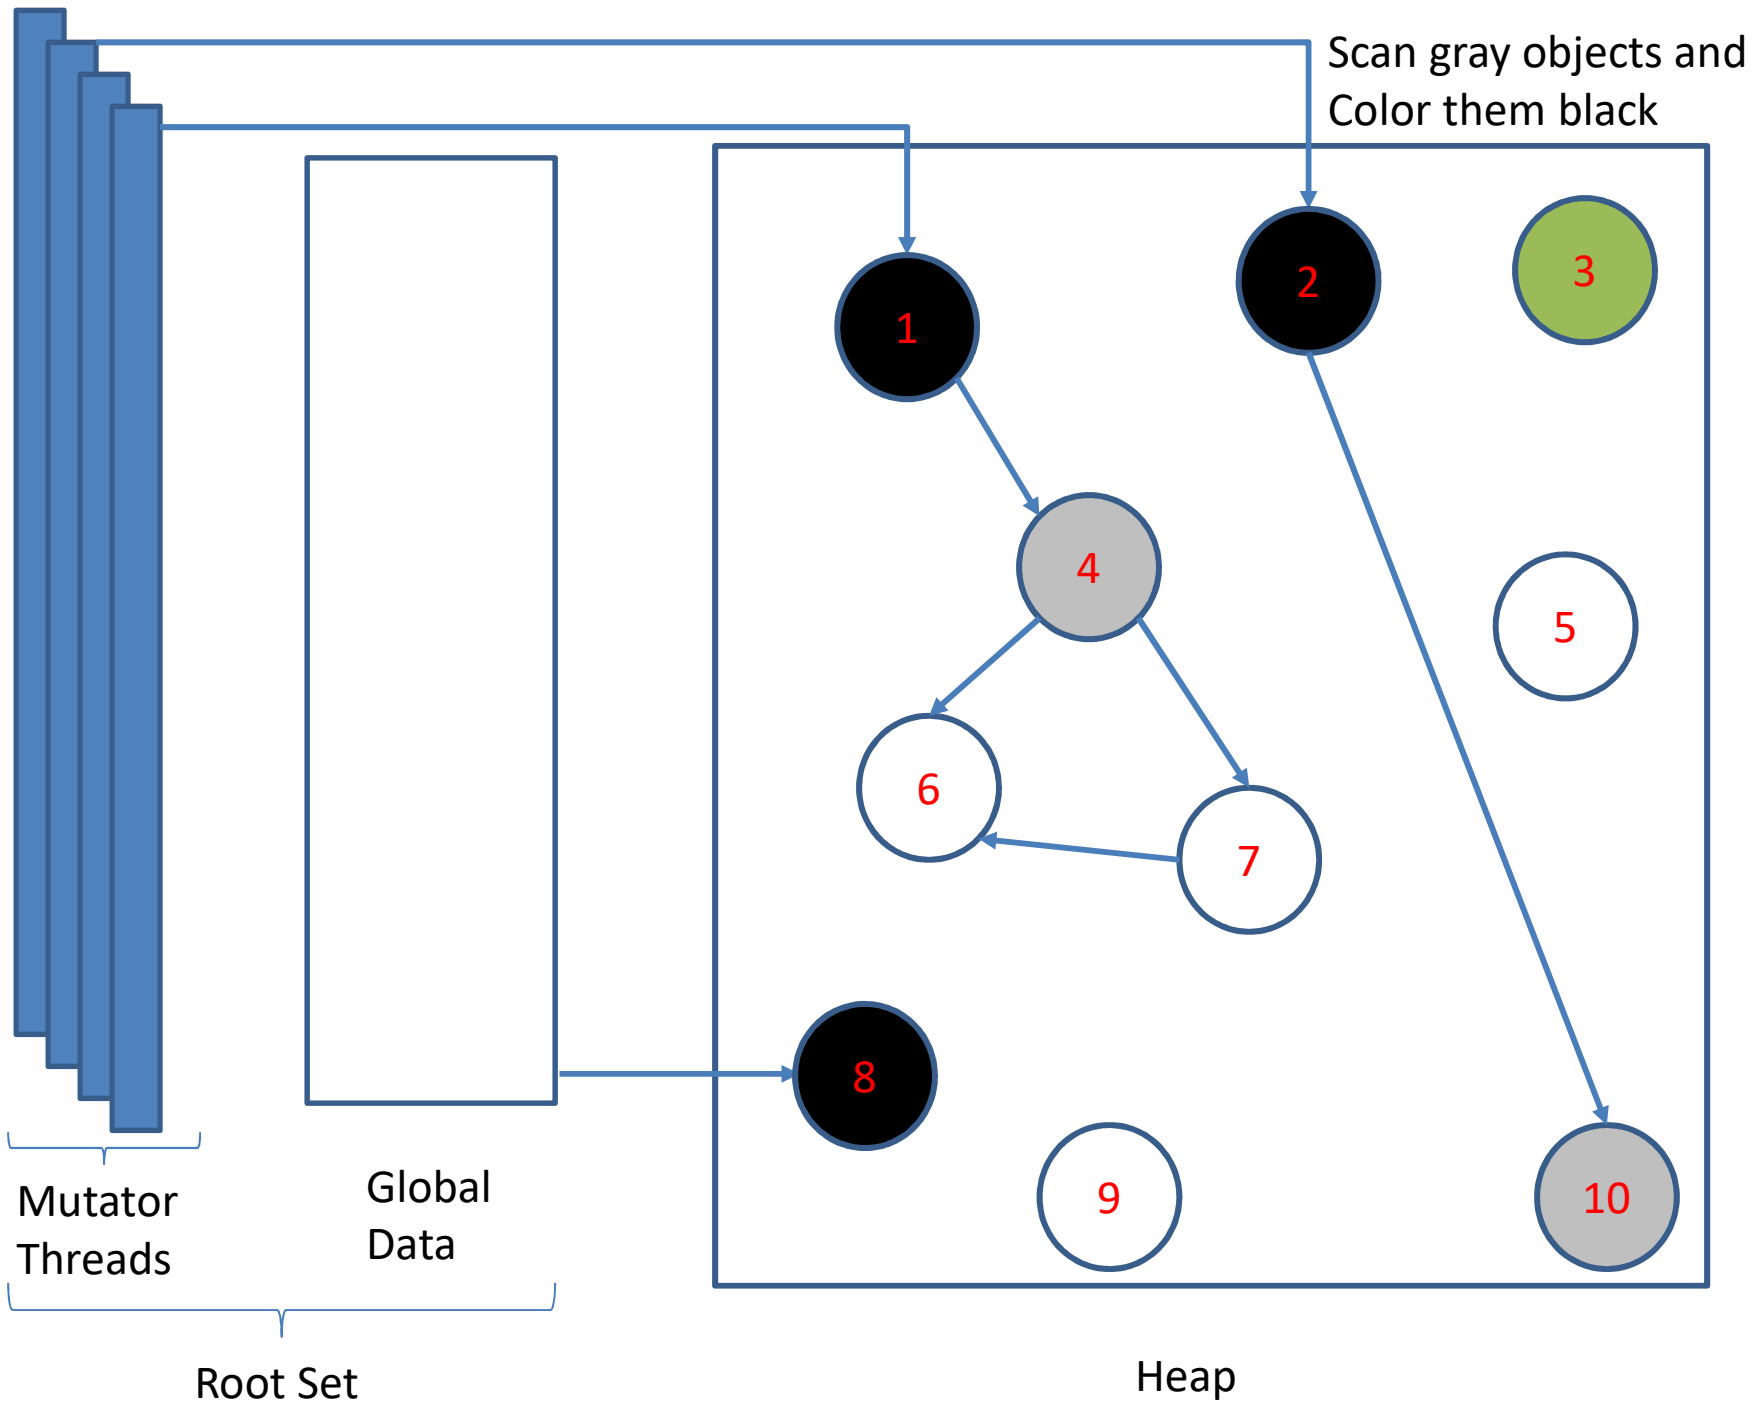


Gray




White






Black




Free


Last

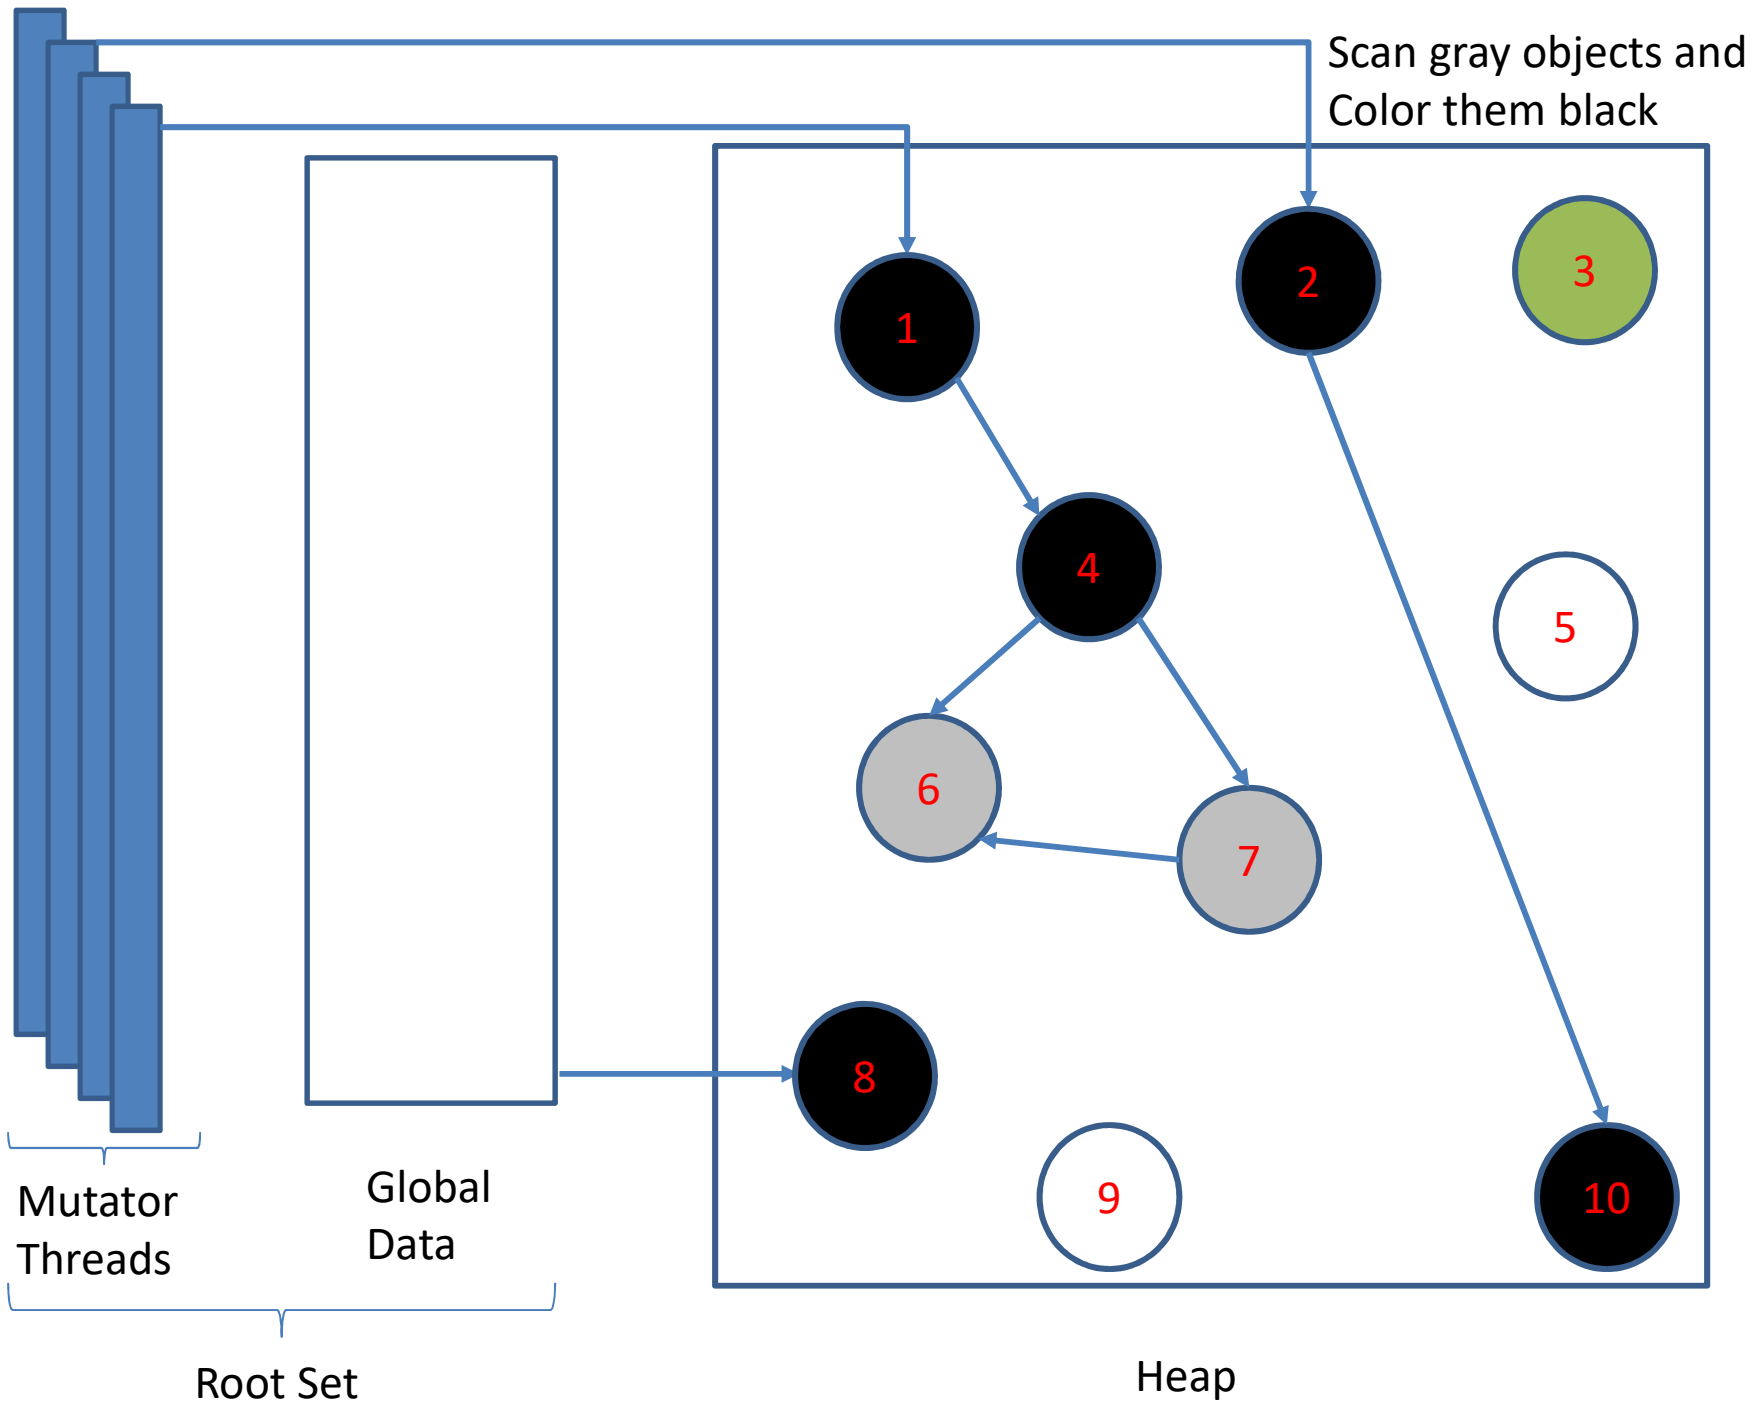


Gray




White






Black




Free


Last






Gray




White



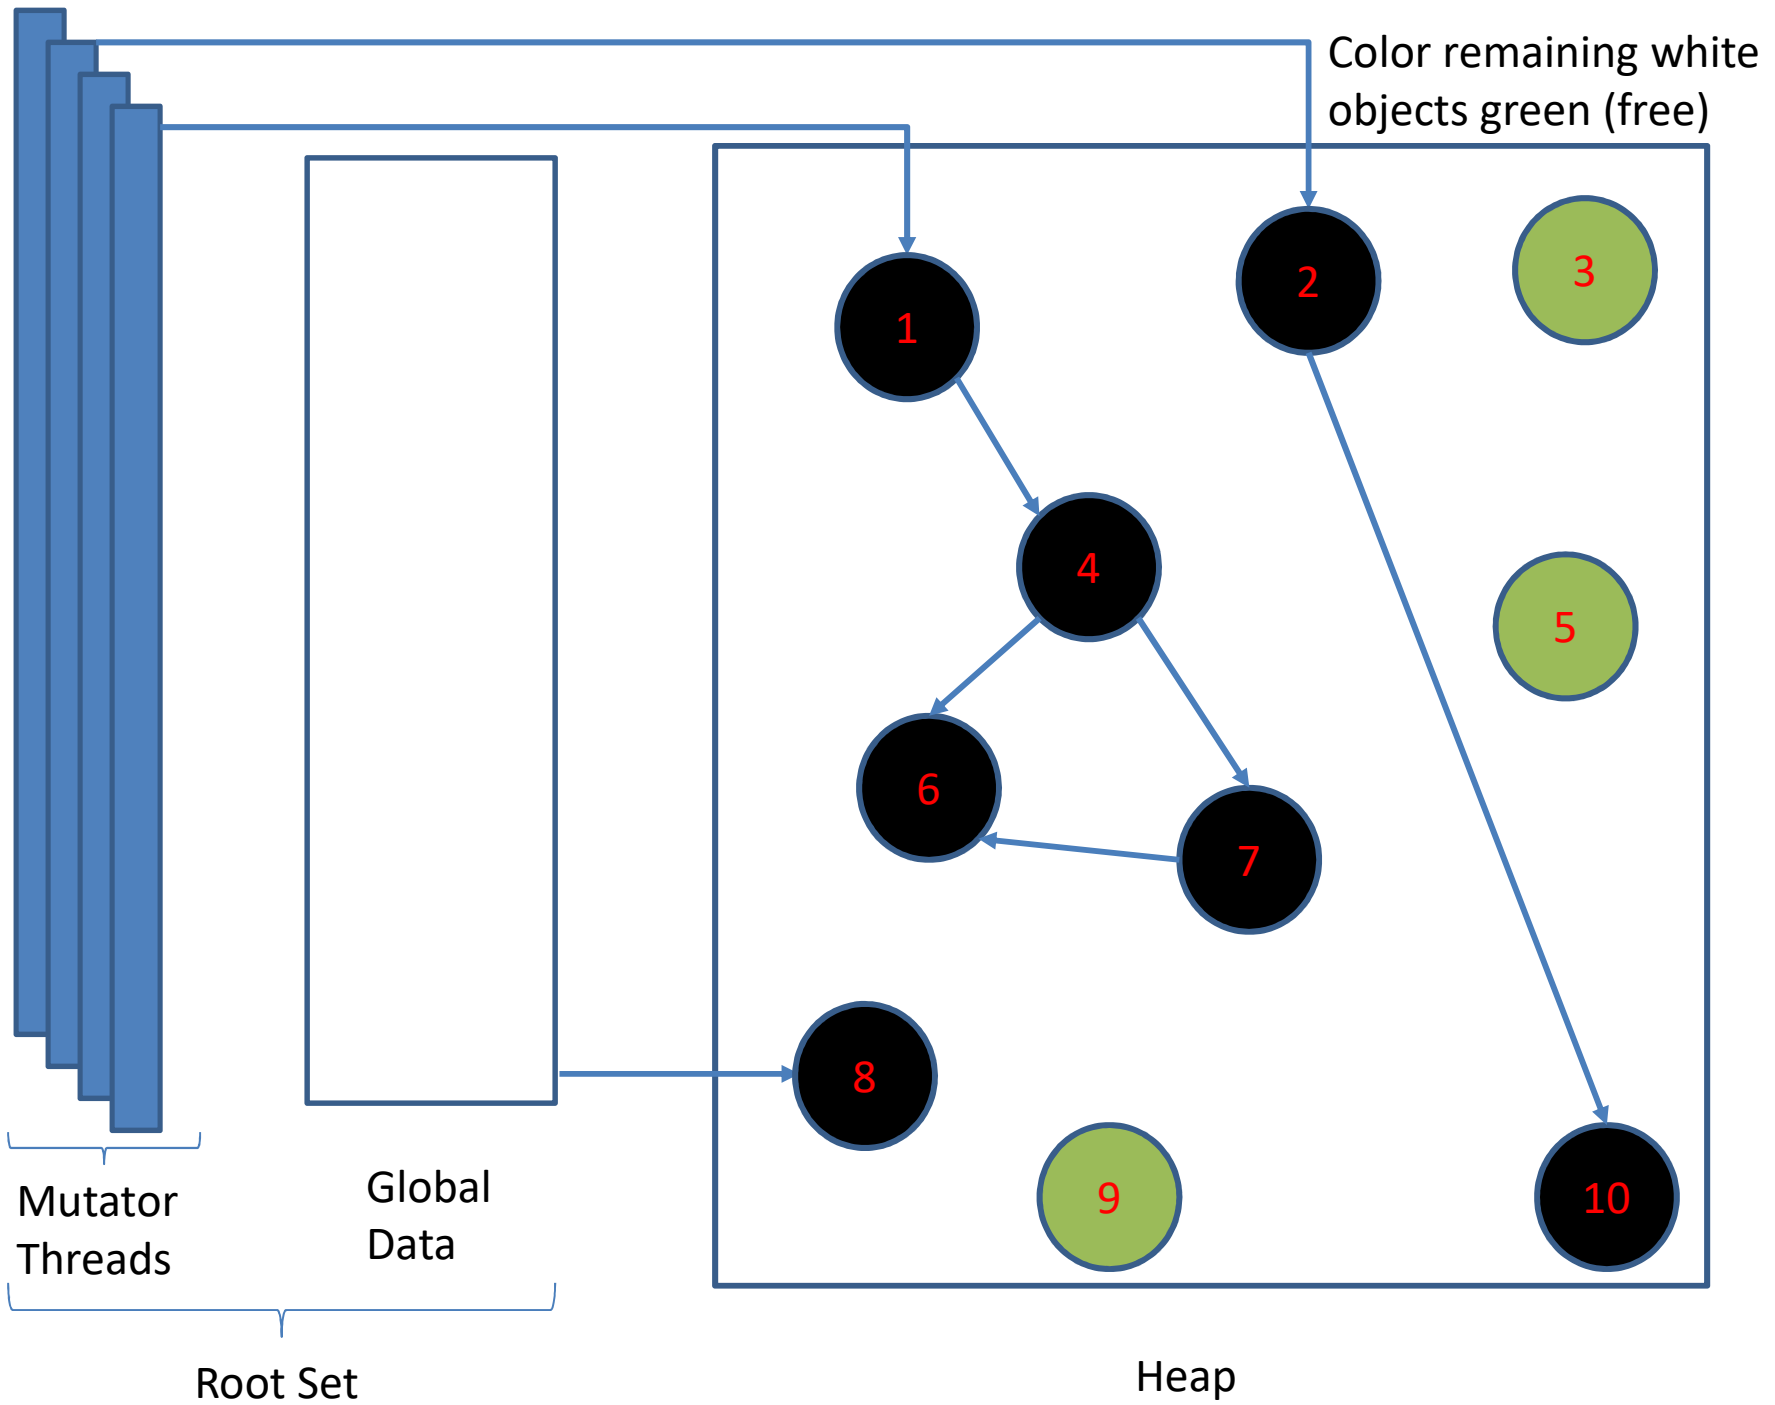





Gray = NULL








Black




Gray = NULL

Free


Last




White = NULL



Black = NULL





Gray = NULL



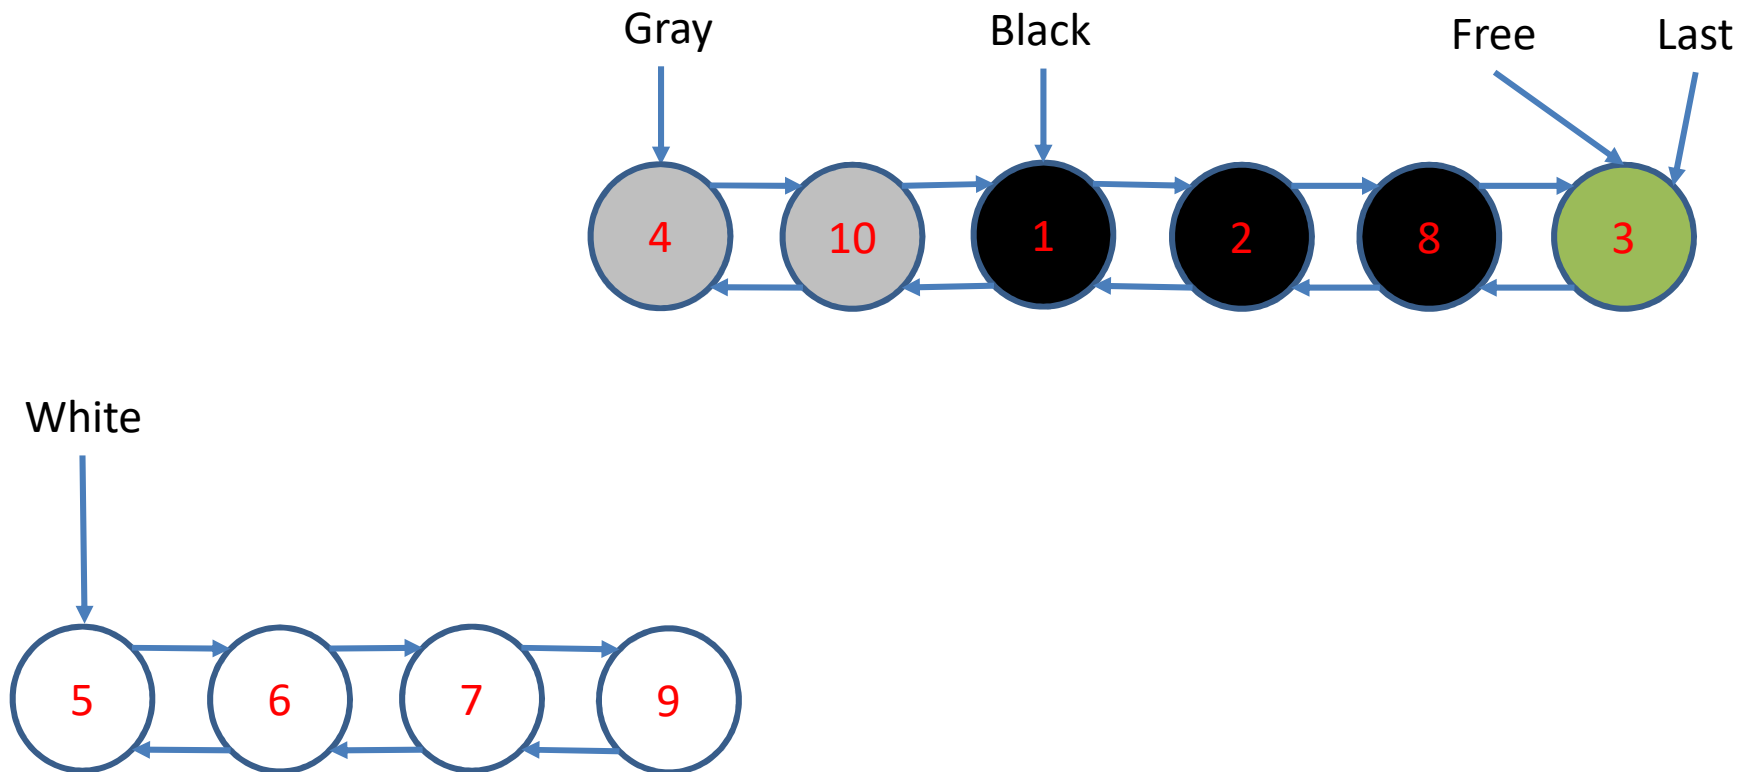
# Let's improve this

- Notice that every live object changes from white, to gray, to black.
- We can make this process more efficient and automatic by connecting some of the linked lists.
- We are also going to color newly allocated objects black.
- We'll examine why later.


Scan/Trace root set



## Tracing object graph




## Tracing object graph






Gc cycle complete  
when Gray == Black




Add garbage to free set




White = NULL

Flip to start new gc cycle

Gray = NULL



White



# Atomic GC Inner loop

```
void full_gc() {  
    flip();  
    scan_root_set();  
  
    do {  
        scan_gray_set();  
    } while (gray != black);  
  
    free_garbage();  
  
    gc_count = gc_count + 1;  
}
```

## Part 2:

# Real-time Concurrent GC

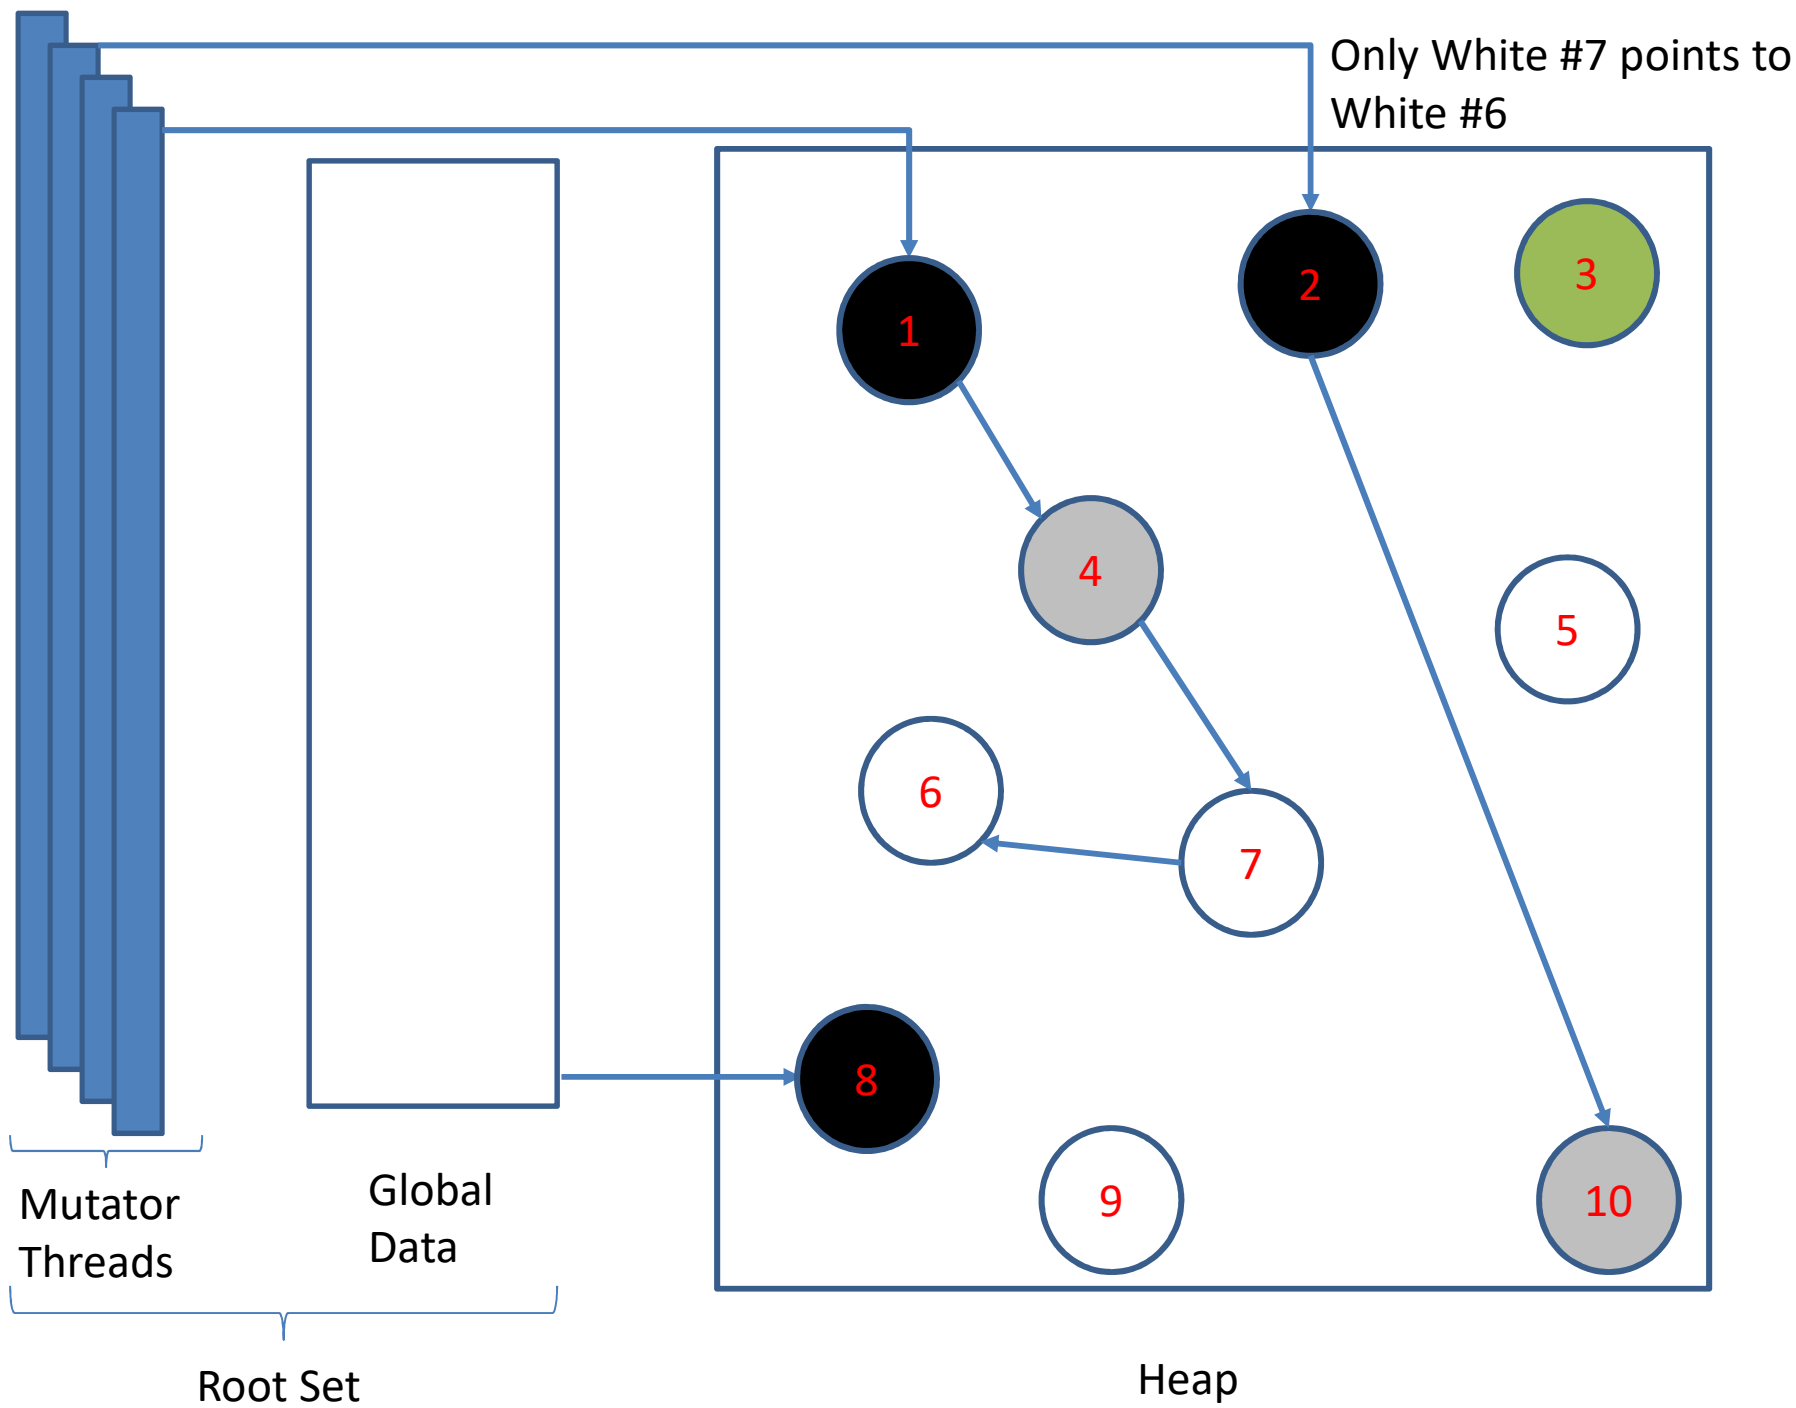
- How do we make the atomic collector a concurrent collector?
  1. Only stop the mutators briefly at the beginning of a new gc cycle when flipping.
  2. Introduce a “write barrier” to allow the mutator and gc threads to run concurrently after that for the rest of the cycle.

# Snap-shot at Beginning

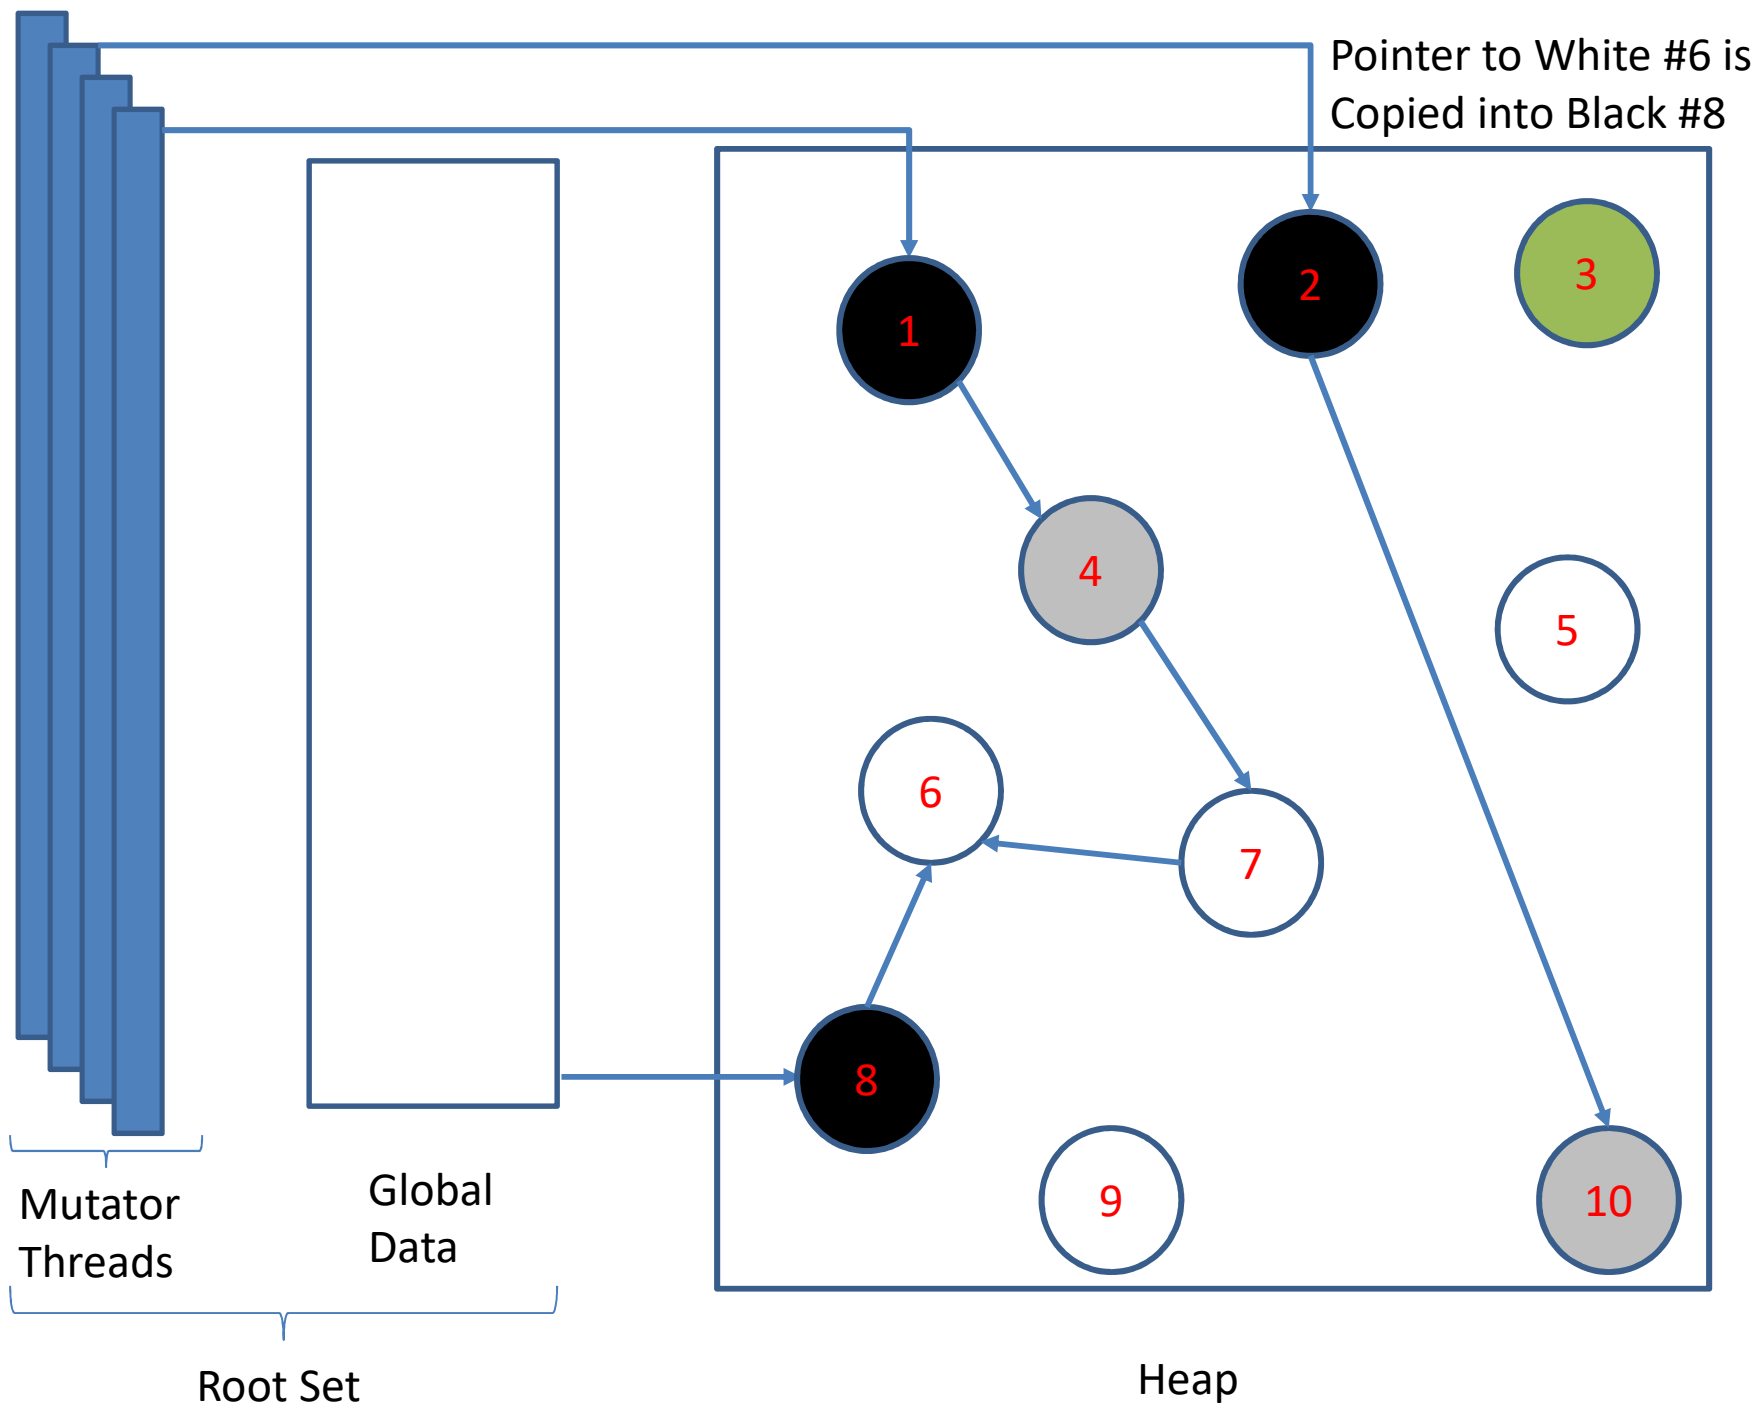
- Conceptually take a complete “snap-shot” of the memory graph when we flipped.
- We are going to retain all objects that are reachable at that time, even if some of them become garbage later during this gc cycle.
- The write barrier allows this by intercepting any mutator pointer writes that might corrupt the snapshot.

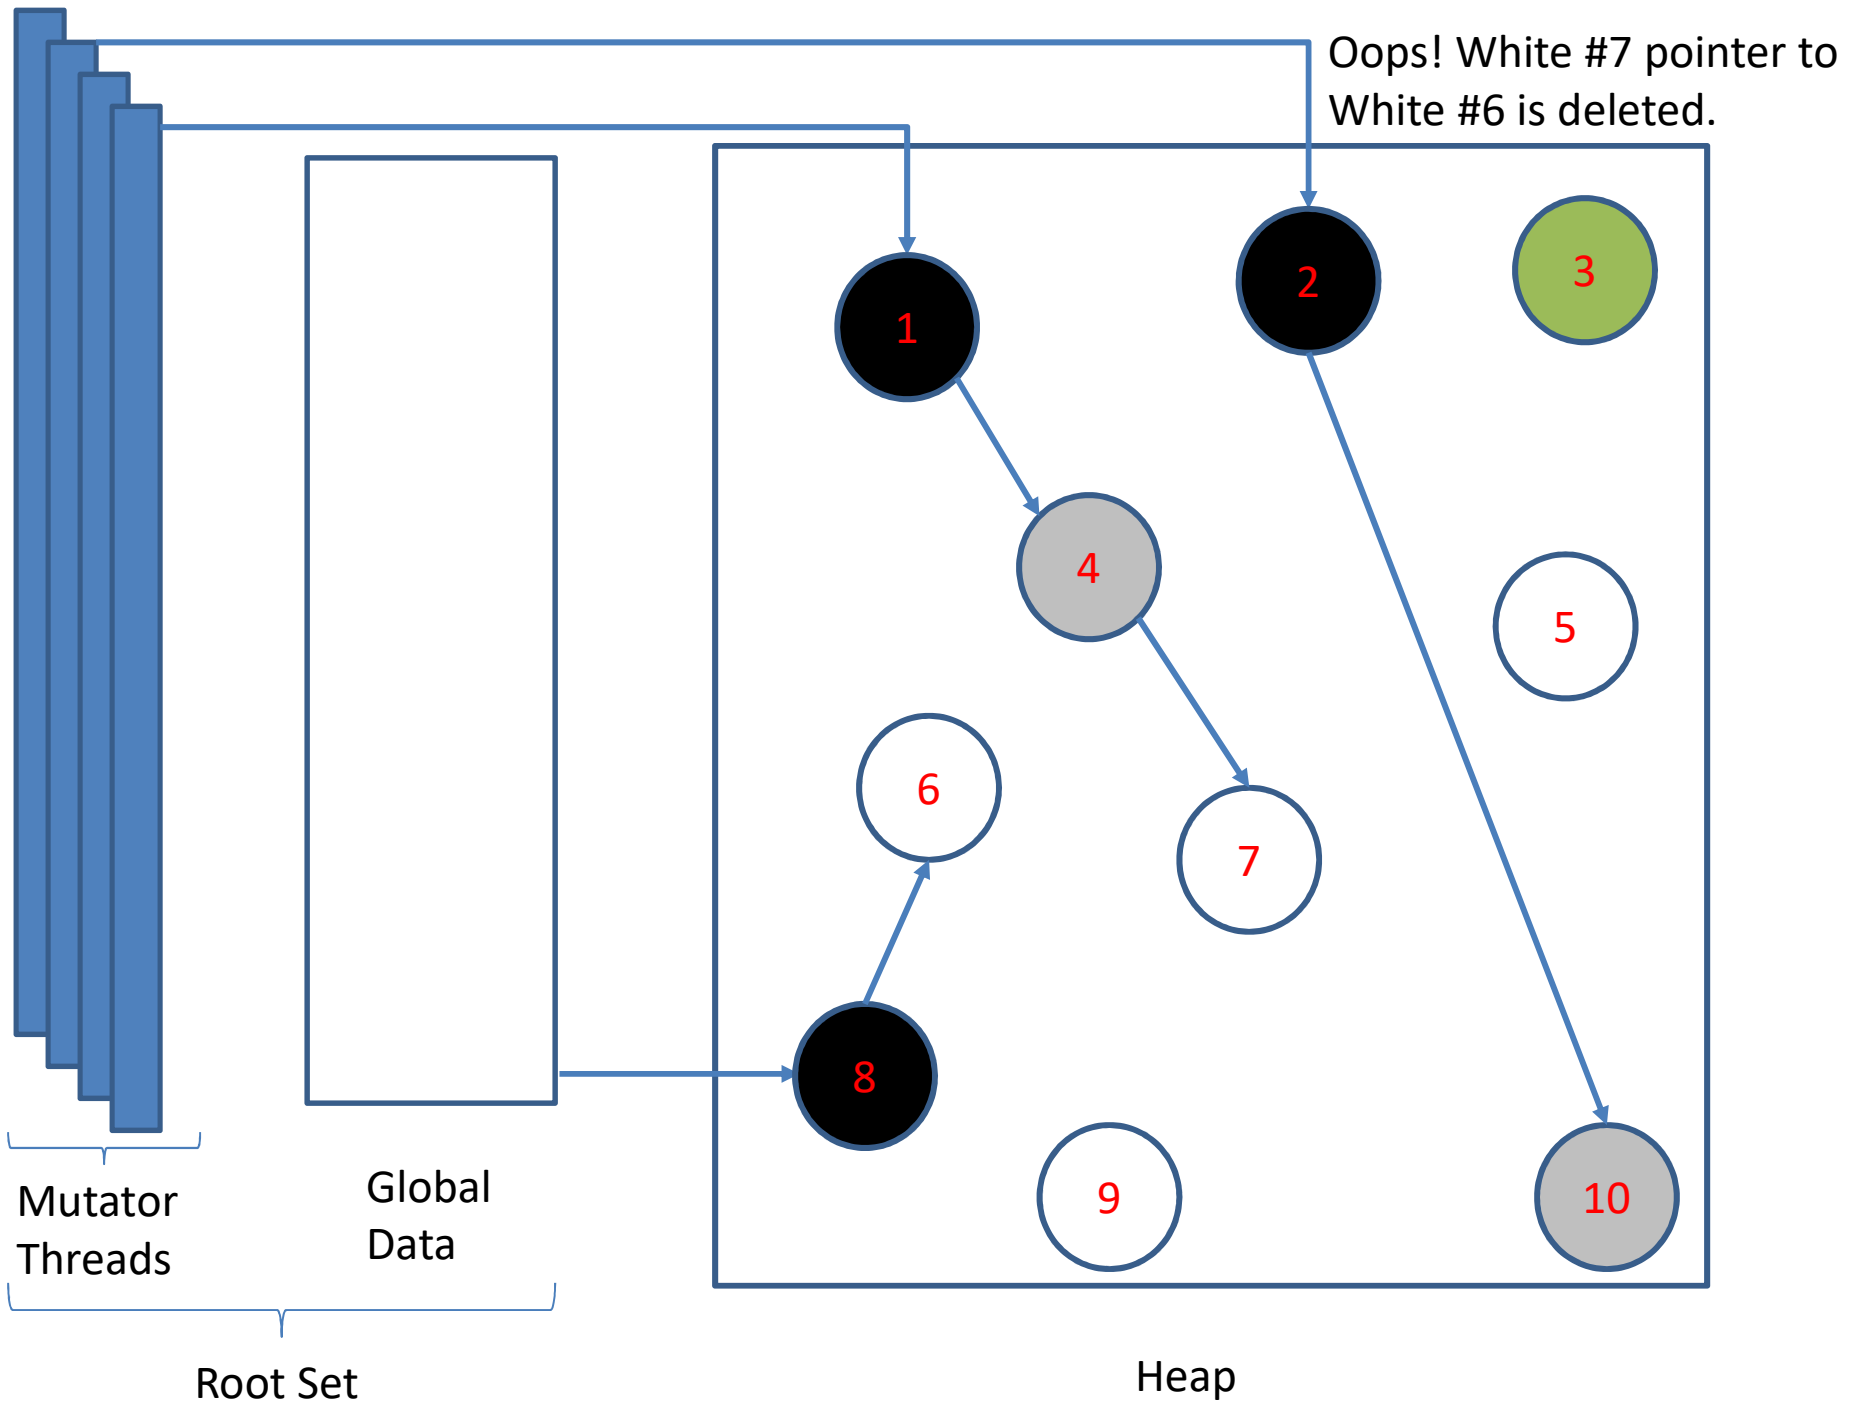
# Why do we need a Snap-shot?

- Suppose a pointer to a white object is stored in exactly one place.
- What if we copy that pointer into another object that is already marked black, and then overwrite the original pointer.
- We have now lost the original pointer, and placed a copy into a black object that will not be scanned again.
- Without a snap-shot, the white object will incorrectly become garbage because it was not traced before getting overwritten.
- With a snap-shot the object stays alive because it was alive when we flipped.





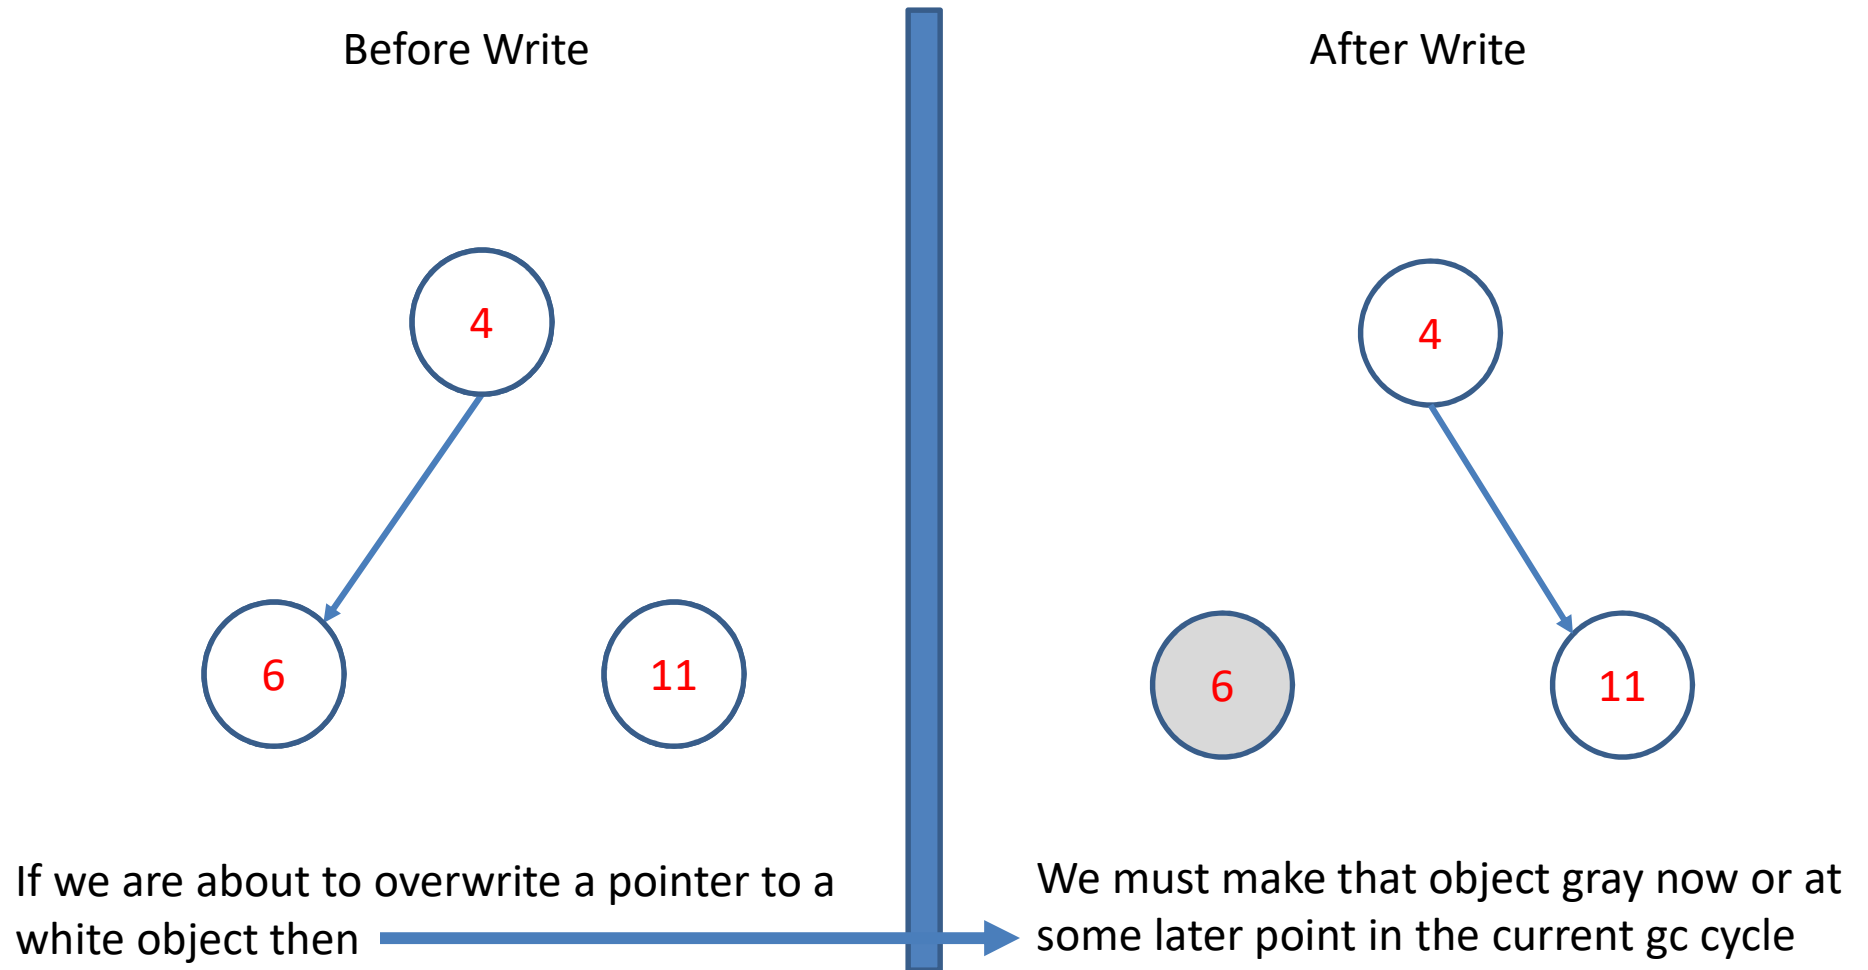




# The Write barrier

- The write barrier intercepts pointer writes and examines the pointer about to be overwritten.
- If a pointer to a white object is about to be overwritten, we must make that object gray now, or at some later point in this gc cycle.
- We don't need to apply this to initializing pointer writes (i.e – Writing pointers into a newly allocated object).
- We'll also see soon why we also don't need to apply this to stack or register writes.

# Write Barrier



# Flipping

- When flipping we pause every mutator thread.
- Each paused mutator copies its stack and register state, and then it may resume.
- We don't need to apply the write barrier to the stack or registers because we atomically copy them somewhere when we flip.
- Each mutator does this independently and only sees a pause time proportional to its own stack and register copy time.
- Flipping completes after every mutator has resumed. This isn't necessary, but it's simple and usually quick.

# Does this really work? - Induction

- Prove: All objects alive at the start of a gc cycle will be retained at the end of the cycle.
- Base case: All objects directly reachable from the saved stack and register state after a flip will be retained after tracing.
- Induction step: Tracing the transitive closure of the object graph reachable from the base case, in conjunction with a mutator write barrier, guarantees that all objects alive at the start of the gc cycle are retained at the end of the cycle.

# Flipping and the Write Barrier

- We can trade off flip pause time and efficiency.
- If we want no mutator pauses when flipping, we need to apply the write barrier to stack and register pointer writes.
- This costs a lot in performance.
- In exchange, we don't ever need to explicitly stop the mutators and copy the stacks and registers.
- We only need to restrict all mutator allocation for the short time it takes to swap the marked and unmarked color. No explicit pauses or copying are required.
- This probably only makes sense for extremely pause time sensitive applications.

# Conservative versus Perfect Collection

- Perfect collection means we can accurately identify all heap pointers, including pointers from stacks and registers.
- Perfect collection only requires Tri-Color marking (no green).
- Conservative collection loosens this constraint and treats anything that looks like a pointer as a pointer.
- Conservative collection uses Quad-Color marking so conservative pointers don't try to "retain" free objects that might look like garbage.
- Conservative collection is helpful in using gc with languages like C, because we do not have "perfect" metadata for stack and register pointers.
- Providing pointer metadata when allocating helps to produce a practical blend of conservative and perfect collection



# Allocate Black

- We allocate black because new objects were not alive when we flipped, so they are not part of the snap-shot.
- New objects could be incorrectly classified as garbage at the end of the gc cycle if we didn't allocate black.
- Unfortunately, this extends the time it takes to reclaim newly allocated objects that quickly become garbage. They can only be reclaimed in the gc cycle that starts after their allocation.

# Continuously Concurrent GC

- There is no Oracle to tell you when an object becomes garbage.
- A “maximally efficient” tracing gc would collect garbage the moment it is created.
- If you have cores (and power) to burn, then continuously gc on as many threads as make sense for your application.
- This can dramatically reduce the maximum heap size.
- Trade off is faster gc vs larger heaps.
- Doing gc work only in response to allocation work is often a mistake. See first bullet point.
- Only slow/stop gc when you have no live objects or you want to save power or use some gc cores for mutator work.

# The Coalescer

- Coalescing is an optional but valuable way to minimize heap size.
- Coalescing competes with allocation on mutator threads to “unallocate” free objects and return whole pages of free objects to their original unallocated state.

# Status

- Single threaded gc runs continuously in a simple test program running multiple mutator threads.
- The WCL implementation of Common Lisp was instrumented with the write barrier and converted to use RTGC. Thousands of self recompiles run while the gc continuously runs concurrently.
- Using the coalescer reduced the WCL self recompile heap size by a factor of 8.

# Next Directions

1. Integrate with a popular language runtime – probably the JDK, but maybe another language.
2. Use RTGC as a replacement for malloc/free in conventional programs C/C++.
3. Use RTGC as a replacement for ARC (Automatic Reference Counting) Apple's Objective-C and Swift language. Reclaim circular memory graphs without burdening developers!
4. Measure and improve RTGC cache friendliness.

Copyright © Wade Lawrence Hennessey 2015 - 2017